

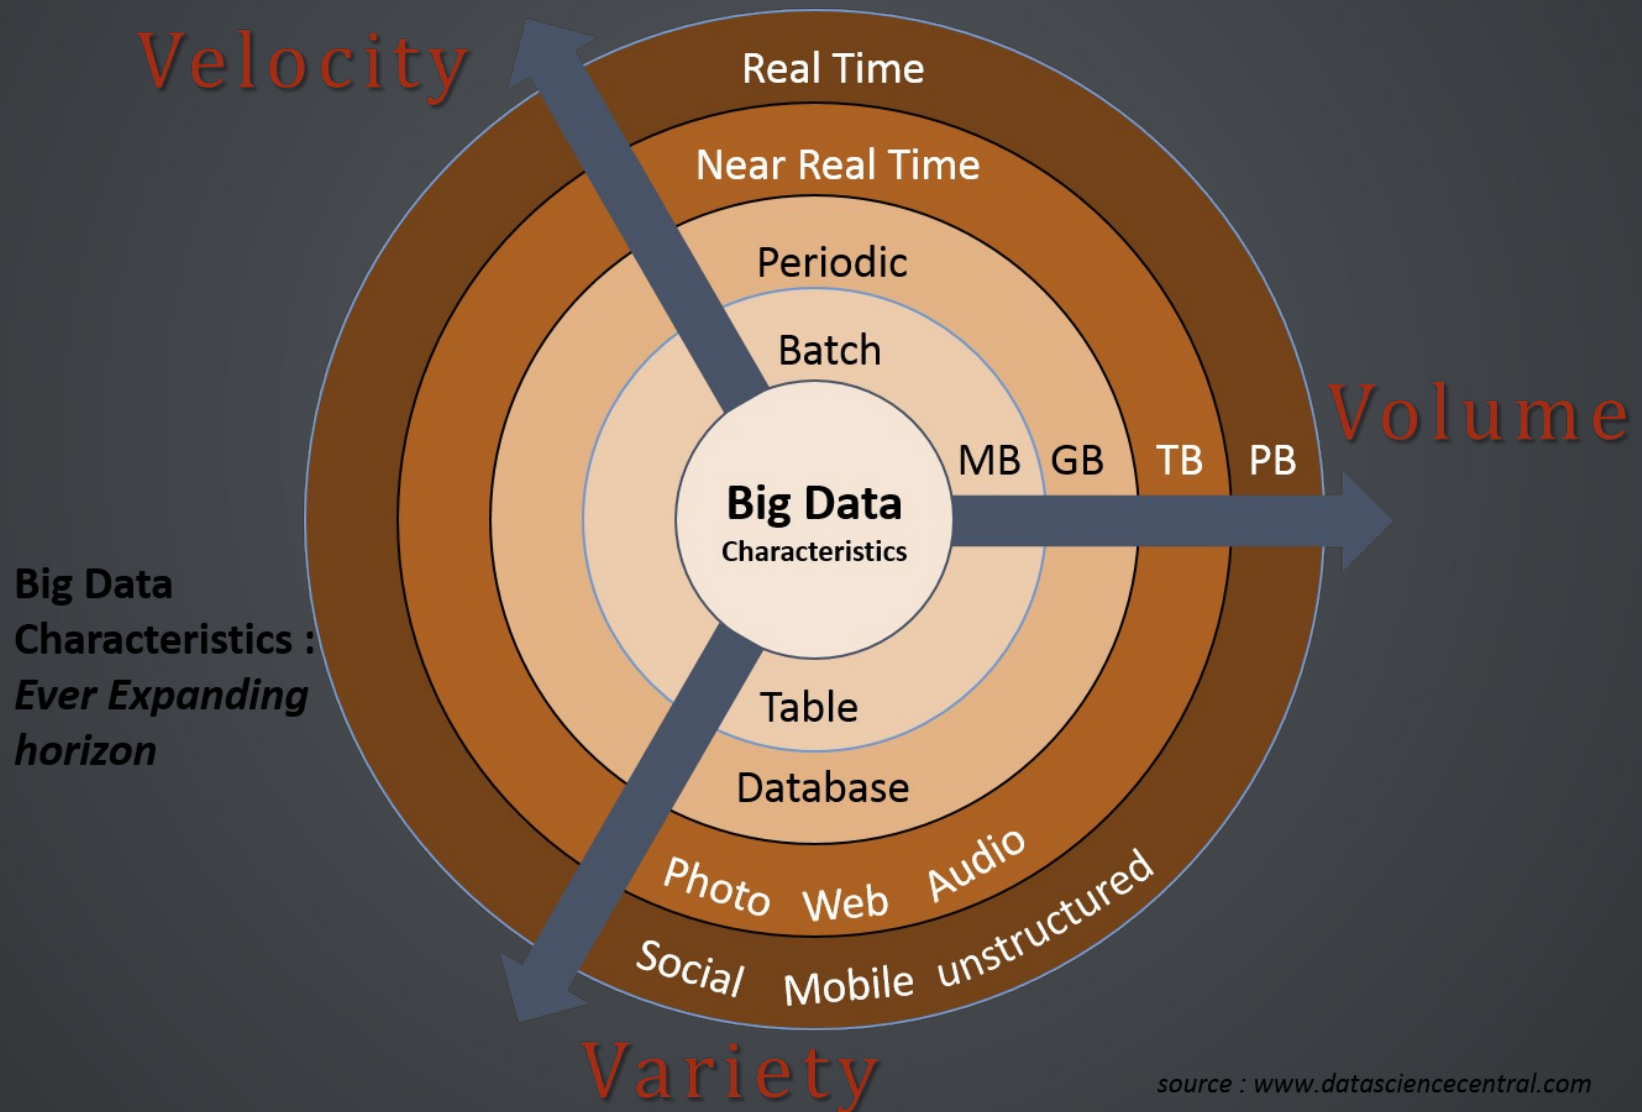
# Services Overview

Javier Cacheiro



# **1. Big Data & Storage**

# 3Vs of Big Data



# DATA CENTRIC

Compute centric:

Bring the **data** to the **computation**

Data centric:

Bring the **computation** to the **data**

# Hardware Infrastructure

- 38 nodes: 4 masters + 34 slaves

Storage capacity **816TB**

Aggregated I/O throughput **30GB/s**

**64GB** RAM per node

**10GbE** connectivity between all nodes

# Hardware Master Nodes

- Model: Lenovo System x3550 M5
- CPU: 2x Intel Xeon E5-2620 v3 @ 2.40GHz
- Cores: 12 (2x6)
- HyperThreading: On (24 threads)
- Total memory: 64GB
- Network: 1x10Gbps + 2x1Gbps

Disks: 8x480GB **SSD** SATA 2.5" MLC G3HS

Controller: M5210 1GB Cache **FastPath**

# Hardware Slave Nodes

- Modelo: Lenovo System x3650 M5
- CPU: 2x Intel Xeon E5-2620 v3 @ 2.40GHz
- Cores: 12 (2x6)
- HyperThreading: On (24 threads)
- Total memory: 64GB
- Network: 1x10Gbps + 2x1Gbps
- Disks: 12x 2TB NL **SATA** 6Gbps 3.5" G2HS
- Controller: N2215 SAS/SATA HBA

# BD|CESGA Software

- Hadoop Platform
- PaaS Platform



# Hadoop Platform

- HDFS
- MapReduce
- Hive
- HBase
- Kafka
- Spark

# PaaS Platform

- MongoDB
- Cassandra
- PostgreSQL
- MariaDB
- etc

# Parallel Storage

- HDFS
- Ceph
- Lustre

# Lustre FT2

- 4xOSS Servers
- 2xMDS Servers
- 8xNetApp E5600 (OSTs):
  - 480x2TB HDD SAS 7200 RPM
- 1xNetApp E2724 (MDS):
  - 24x1TB HDD SAS 10.500 RPM

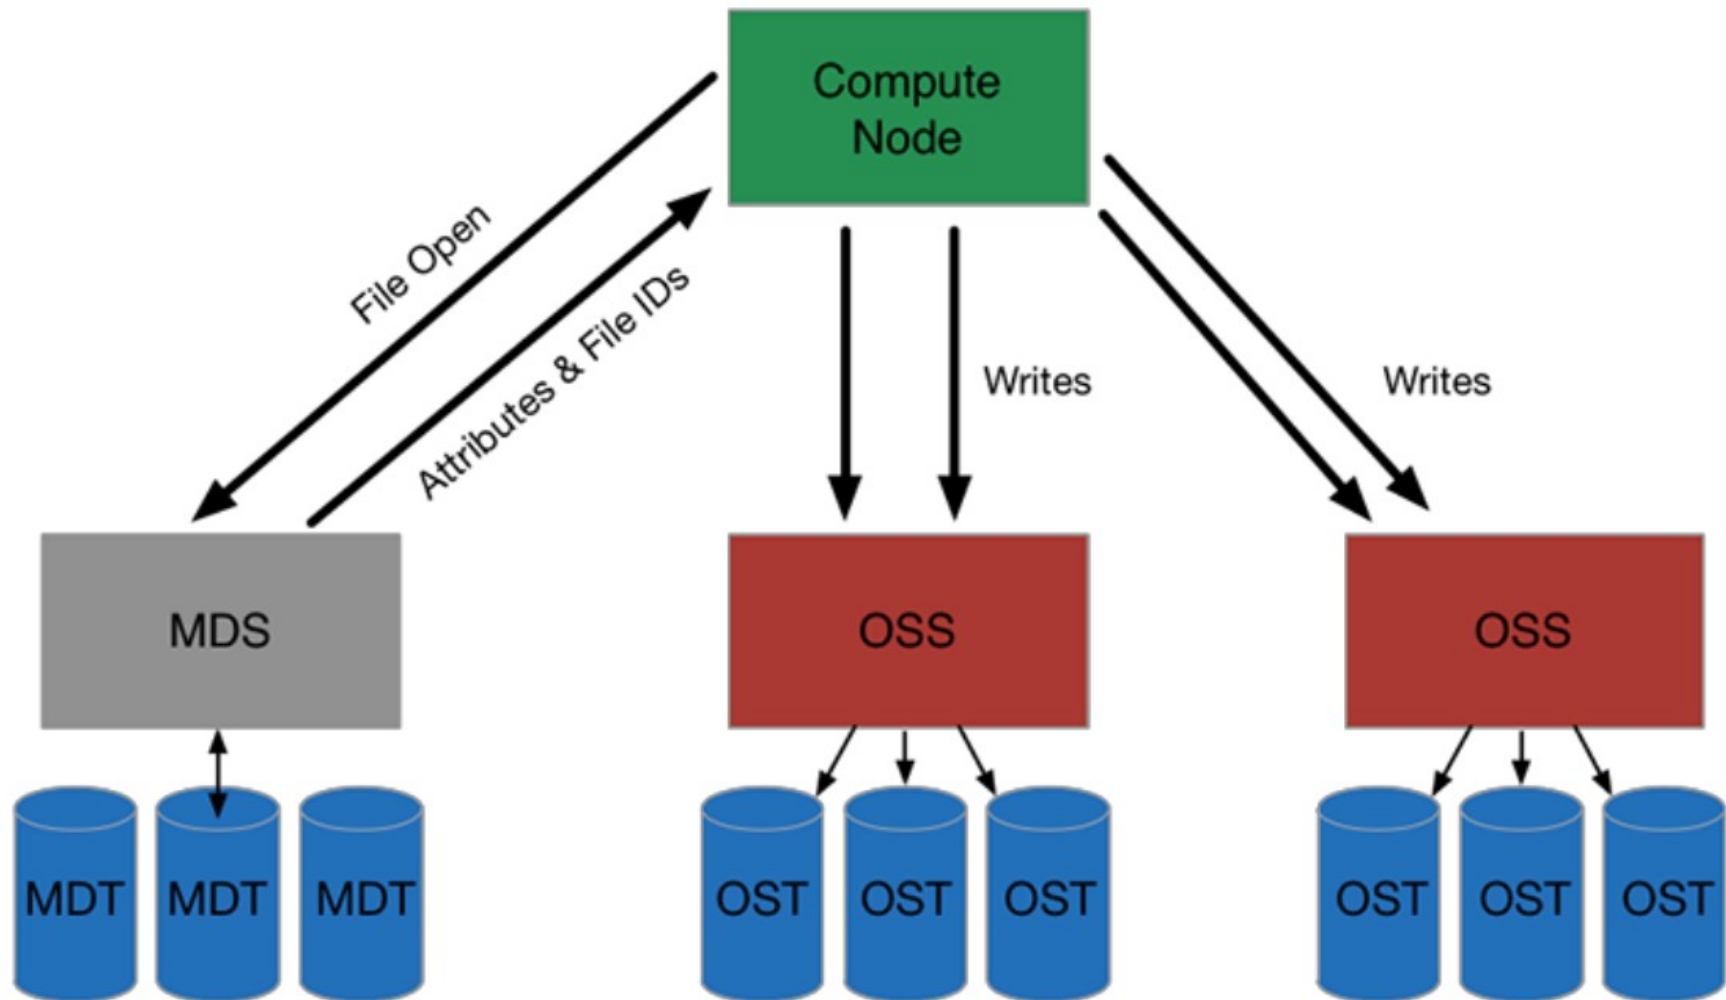
# Lustre FT2


Bandwidth: **24GB/s**

Capacity: **695TB**

Connectivity: **Infiniband FDR 56Gbps**

# Lustre





# FT3 Storage

**Almacenamiento de altas prestaciones**  
**Lustre 1 PB**

**Almacenamiento de alta capacidad**

**Sistema de almacenamiento  
permanente de datos en disco**  
**4 PB Lustre + HSM (Robinhood +  
QSTAR)**

**Sistema de almacenamiento  
permanente en cinta**  
**20 PB**

**Librería de cintas**  
**12x LT09 + 4x LT06**  
**1600 slots**



# Backups: Tape Library



# Netapp FAS9000HA

Big Data/FT2/FT3 HOME

- High availability
- High cost
- Lower parallel performance

# Netapp

Bandwidth: **14GB/s**

Capacity: > **1PB**

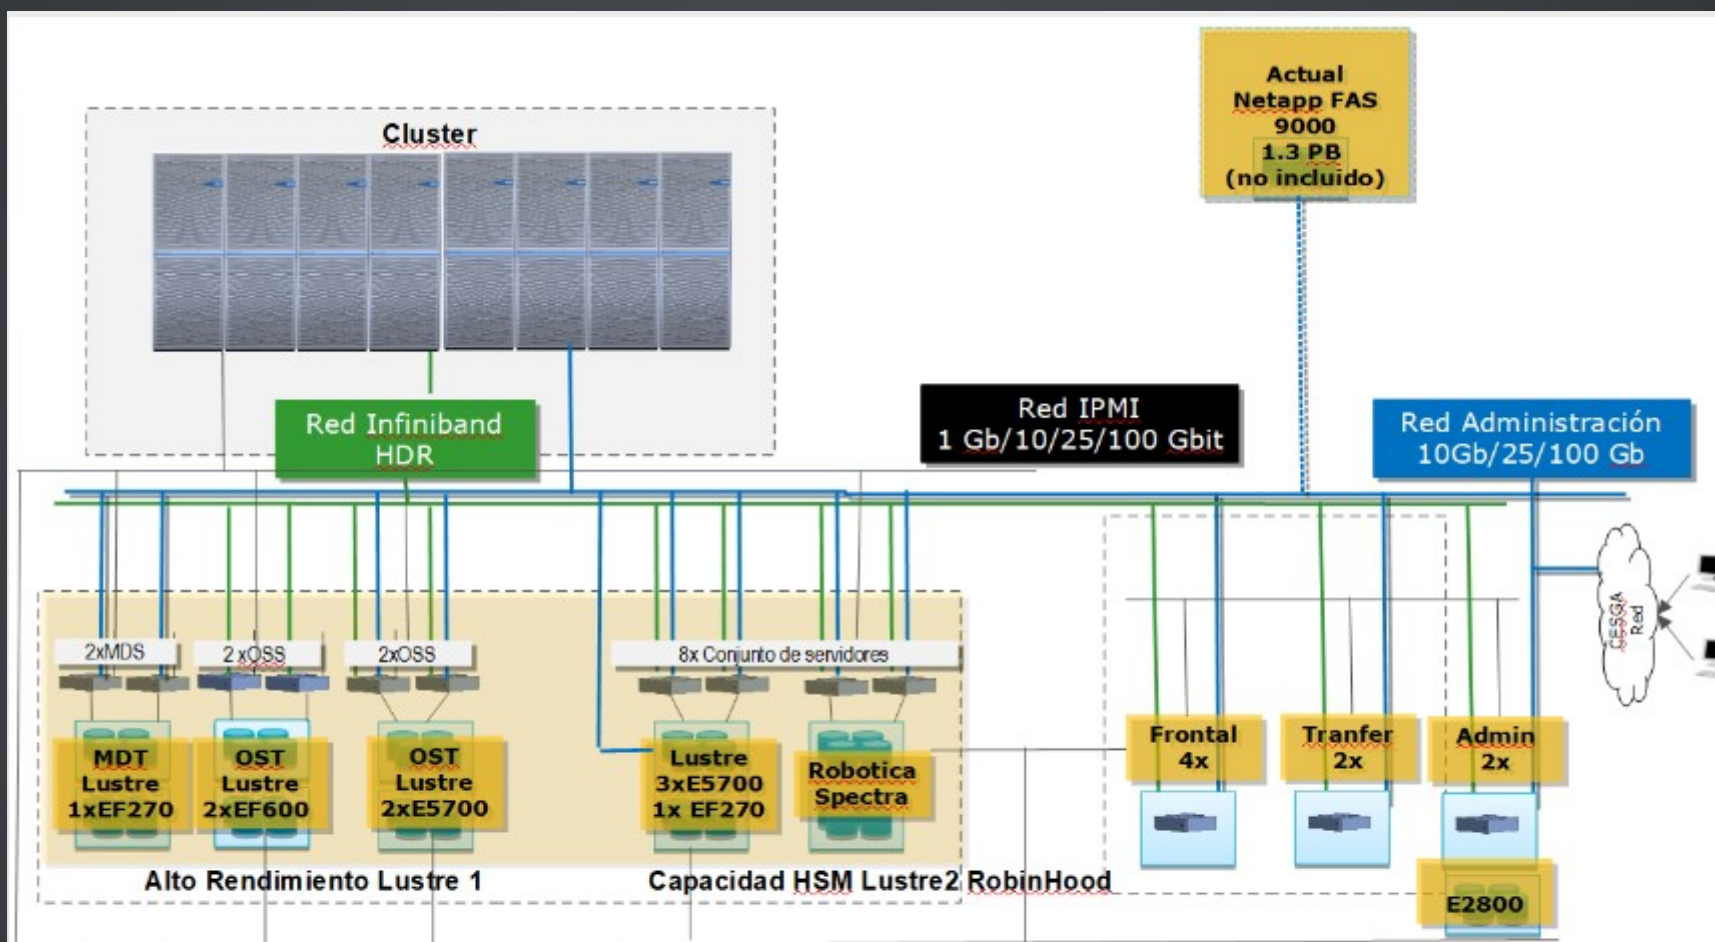
Performance: **1.5M IOPS**



## **2. Computing**

# FinisTerra III (FT3)

- 708 x procesadores Intel Xeon Ice Lake 8352Y de 32 cores, 2.2GHz
- 22.656 núcleos de proceso (cores)
- 144 aceleradores matemáticos GPU Nvidia A100 y Nvidia T4
- 118 TB de memoria agregada
- 355 TB capacidad agregada en discos SSD NVMe
- Red Infiniband HDR 100



**Potencia de cómputo de 4 PetaFLOPS**



	Thin nodes	GPU nodes	Viz nodes	Fat nodes	Transfer nodes	QLM node
Cantidad	256	64	16	16	2	Simulador de circuitos cuánticos Atos 30 qbits
Procesador	2x Intel Xeon Ice Lake 8352Y 32 cores 2.2 GHz					
Cores	64					
Memoria	256 GB			2 TB	256 GB	
Disco local	960 GB SSD NVMe			1920 GB SSD NVMe	960 GB SSD NVMe	
GPU	2x Nvidia A100		Nvidia T4			
Red baja latencia	Infiband HDR 100					



# **3. Cloud**

	Cloud 1	Cloud 2	Cloud 3
Year	2009	2013	2021
Cores	1.200	1.464	6.240
Memory	2,4 TB	3,5 TB	24 TB
Storage	5,6 TB	63 TB	3.000 TB
Software	Opennebula 1	OpenNebula 5	OpenStack Xena

